

Walkthrough: intro

In this binary, we are going to take a look at the basic toolkit that we use to disassemble and analyze binaries. This binary isn't meant to be difficult but rather as an introduction to assembly.

First Checks

When you download the binary, the first thing I recommend is to run it. Because you downloaded the binary from the internet, execution will be turned off by default. To fix this:

```
$ chmod +x intro
$ ./intro
```

We are provided the following output:

```
$ ./intro
Enter the flag here: flag
Nope, try again!
```

It seems that we need to just enter the flag and it will tell us if we get it right.

Second, we check the security measures of the file. Use `checksec` for this. If your program says `checksec: not found`, let me know and I think I know the problem. `checksec intro` prints the following:

```
[*] '/home/joybuzzer/Documents/vunrotc/live/00-introduction/intro/intro'
Arch:      amd64-64-little
RELRO:     Full RELRO
Stack:     No canary found
NX:        NX enabled
PIE:       PIE enabled
```

Let's break this down:

- **Arch: amd64-64-little** - This means that the binary is compiled for a 64-bit architecture. This is important because it means that we will be using 64-bit registers and instructions.
- **RELRO: Full RELRO** - This means that the Global Offset Table (GOT) is read-only. This is important because it means that we cannot overwrite the GOT to redirect execution to a different function. *More on this later.*
- **Stack: No canary found** - This means that the stack is not protected by a Canary. This is important because it means we can overwrite stack data without the program knowing.
- **NX: NX enabled** - This means that the stack is not executable. This is important because it means we cannot put outside instructions into the program and force the program to execute them.

- **PIE: PIE enabled** - This means that the binary is compiled with Position Independent Executable (PIE) enabled. This is important because it means that the binary will be loaded into a random location in memory. *More on this later.*

What does all this mean?

- We are dealing with a 64-bit binary.
- If we wanted to write to a specific place in memory, we can't because the stack & heap will be in different places every execution.

Let's start disassembling this binary.

Disassembly

Running this binary with **gdb** will let us see the assembly behind the program. To do this, we run **gdb intro**. From here, the best thing to check are the available functions (**info functions**):

```
gef> info functions
All defined functions:

Non-debugging symbols:
0x0000000000001000 _init
0x0000000000001080 __cxa_finalize@plt
0x0000000000001090 puts@plt
0x00000000000010a0 printf@plt
0x00000000000010b0 fgets@plt
0x00000000000010c0 strcmp@plt
0x00000000000010d0 malloc@plt
0x00000000000010e0 _start
0x0000000000001110 deregister_tm_clones
0x0000000000001140 register_tm_clones
0x0000000000001180 __do_global_dtors_aux
0x00000000000011c0 frame_dummy
0x00000000000011c9 main
0x0000000000001258 _fini
```

Most of these functions are standard library functions. The ones suffixed **@plt** are external library functions. The only function that we are interested in is **main**. Let's take a look at the assembly for **main** (**disas main**):

```
gef> disas main
Dump of assembler code for function main:
0x00000000000011c9 <+0>: endbr64
0x00000000000011cd <+4>: push    rbp
0x00000000000011ce <+5>: mov     rbp, rsp
0x00000000000011d1 <+8>: sub     rsp, 0x20
0x00000000000011d5 <+12>: mov     DWORD PTR [rbp-0x14], edi
0x00000000000011d8 <+15>: mov     QWORD PTR [rbp-0x20], rsi
0x00000000000011dc <+19>: mov     edi, 0x20
0x00000000000011e1 <+24>: call    0x10d0 <malloc@plt>
```

```

0x000000000000011e6 <+29>:  mov     QWORD PTR [rbp-0x8],rax
0x000000000000011ea <+33>:  lea     rax,[rip+0xe13]          # 0x2004
0x000000000000011f1 <+40>:  mov     rdi,rax
0x000000000000011f4 <+43>:  mov     eax,0x0
0x000000000000011f9 <+48>:  call    0x10a0 <printf@plt>
0x000000000000011fe <+53>:  mov     rdx,QWORD PTR [rip+0x2e0b]      #
0x4010 <stdin@GLIBC_2.2.5>
0x00000000000001205 <+60>:  mov     rax,QWORD PTR [rbp-0x8]
0x00000000000001209 <+64>:  mov     esi,0x20
0x0000000000000120e <+69>:  mov     rdi,rax
0x00000000000001211 <+72>:  call    0x10b0 <fgets@plt>
0x00000000000001216 <+77>:  mov     rax,QWORD PTR [rbp-0x8]
0x0000000000000121a <+81>:  lea     rdx,[rip+0xdf9]          # 0x201a
0x00000000000001221 <+88>:  mov     rsi,rdx
0x00000000000001224 <+91>:  mov     rdi,rax
0x00000000000001227 <+94>:  call    0x10c0 <strcmp@plt>
0x0000000000000122c <+99>:  test    eax,eax
0x0000000000000122e <+101>: jne     0x1241 <main+120>
0x00000000000001230 <+103>: lea     rax,[rip+0xdfd]          # 0x2034
0x00000000000001237 <+110>: mov     rdi,rax
0x0000000000000123a <+113>: call    0x1090 <puts@plt>
0x0000000000000123f <+118>: jmp     0x1250 <main+135>
0x00000000000001241 <+120>: lea     rax,[rip+0xe03]          # 0x204b
0x00000000000001248 <+127>: mov     rdi,rax
0x0000000000000124b <+130>: call    0x1090 <puts@plt>
0x00000000000001250 <+135>: mov     eax,0x0
0x00000000000001255 <+140>: leave
0x00000000000001256 <+141>: ret
End of assembler dump.

```

Now this is a lot. The way I recommend to do this is to go from external function call to external function call, understanding what gets passed to them. Let's do this one at a time.

Hold on, One more thing

We need to understand how 64-bit functions pass parameters. These are done **via registers**. The first 6 parameters are passed in the following registers:

- **rdi** (sometimes as **edi**, the lower 4-bytes of **rdi**)
- **rsi** (sometimes as **esi**, the lower 4-bytes of **rsi**)
- **rdx** (sometimes as **edx**, the lower 4-bytes of **rdx**)
- **rcx** (sometimes as **ecx**, the lower 4-bytes of **rcx**)
- **r8** (sometimes as **r8d**, the lower 4-bytes of **r8**)
- **r9** (sometimes as **r9d**, the lower 4-bytes of **r9**)

The difference between **rdi** and **edi** is confusing. **rdi** and **edi** are connected to the same storage location, meaning that changing **edi** affects **rdi**, and vice versa. **edi** is just the bottom half of **rdi**, as we see below.

```
rdi = [ _ _ _ _ _ _ _ _ ]
      edi = [ _ _ _ _ ]
```

Assembly often passes parameters, especially parameters with small values, using the lower four bytes (i.e. `edi`). Assembly does this because *it's faster*, but is no different than passing all 8 bytes. The same goes for `rsi`, `rdx`, `rcx`, `r8`, and `r9`.

Now let's continue.

malloc

`malloc()` in C is how we allocate memory. It makes space on the heap for us to use. If you want to see the parameters for `malloc()`, we look at the [man](#) pages. This is where the function header is defined. [man malloc](#) shows the following:

SYNOPSIS

```
#include <stdlib.h>
```

```
void *malloc(size_t size);
```

DESCRIPTION

The `malloc()` function allocates `size` bytes and returns a pointer to the allocated memory. The memory is not initialized. If `size` is 0, then `malloc()` returns either `NULL`, or a unique pointer value that can later be successfully passed to `free()`.

This tells us that `malloc()` takes a parameter being the number of bytes that gets passed to it. Since we know that `rdi` is the register that **always** holds the first parameter, we check above the function call for `rdi` being stored.

```
0x000000000000011dc <+19>:  mov     edi,0x20
0x000000000000011e1 <+24>:  call    0x10d0 <malloc@plt>
0x000000000000011e6 <+29>:  mov     QWORD PTR [rbp-0x8],rax
```

Here we see that `edi` (aka `rdi`) is being loaded with `0x20` (32). This means that the equivalent function here is `malloc(0x20)`, which allocates 32 bytes on the heap.

The `str` variable resides in memory at address `rpb-0x8`. Immediately after the `malloc` returns, register `rax` contains the `malloc` return value, and that value is stored in the memory at address `rpb-0x8` (i.e. it is stored in `str`).

printf

Let's check the arguments for `printf`.

SYNOPSIS

```
printf FORMAT [ARGUMENT]...
printf OPTION
```

DESCRIPTION

Print ARGUMENT(s) according to FORMAT, or execute according to OPTION:

The `man` pages aren't super helpful here, but `printf` just prints text to the screen. `rdi` will have the string in memory that we want to print.

```
0x000000000000011ea <+33>: lea    rax,[rip+0xe13]      # 0x2004
0x000000000000011f1 <+40>: mov    rdi,rax
0x000000000000011f4 <+43>: mov    eax,0x0
0x000000000000011f9 <+48>: call   0x10a0 <printf@plt>
```

The first command is the most confusing. `lea` stands for **Load Effective Address**, which takes the address of the memory location and stores it in the register. This is the same as `mov rax, 0x2004`. The next command moves the address of the string into `rdi`, and the last command calls `printf`. *Why does assembly use `lea`?* `lea` is faster than `mov` because it doesn't have to access memory to get the value. It just gets the address and stores it in the register.

After this, it moves this address into `rdi`. If we check what's at this address, we see the following:

```
gef> x/s 0x2004
0x2004: "Enter the flag here: "
```

This makes sense with what we saw in the program output.

fgets

Let's check the `man` pages to see what `fgets` does.

SYNOPSIS

```
char *fgets(char *s, int size, FILE *stream);
```

DESCRIPTION

`fgets()` reads in at most one less than `size` characters from `stream` and stores them into the buffer pointed to by `s`. Reading stops after an EOF or a newline.

If a newline is read, it is stored into the buffer. A terminating null byte (`'\0'`) is stored after the last character in the buffer.

This tells that `fgets()` is an input function. Since `fgets()` takes three arguments, we need to find what `rdi`, `rsi`, and `rdx` are to understand what's going into the function.

```
0x000000000000011fe <+53>:    mov     rdx,QWORD PTR [rip+0x2e0b]    #  
0x4010 <stdin@GLIBC_2.2.5>  
0x00000000000001205 <+60>:    mov     rax,QWORD PTR [rbp-0x8]  
0x00000000000001209 <+64>:    mov     esi,0x20  
0x0000000000000120e <+69>:    mov     rdi,rax  
0x00000000000001211 <+72>:    call    0x10b0 <fgets@plt>
```

Let's find the parameters.

- We see that `rdi` is loaded with `rax`, which is loaded with the address at `rbp-0x8`. From before, we know that this is where the `malloc` happened.
- `rsi` is loaded with `0x20`, meaning we are writing up to `0x20` bytes.
- `rdx` is loaded with the address of `rip+0x2e0b` (which GDB tells us is at `0x4010`). GDB informs us that this is `stdin`, which is the stream that the data comes from.

From this, we know that the C code looks like this:

```
buffer = malloc(0x20);  
fgets(buffer, 0x20, stdin);
```

What does this mean? From earlier, we know that when `malloc` is called, data is written on the **heap**. This means that we're writing data to the heap, *rather than the stack*.

Why is this important? This means that we can't perform many of the stack overflow techniques we are going to learn. There are a series of heap overflow techniques, but they are relatively out of the scope of the screener.

strcmp

`strcmp`, or *string compare*, is how C compares strings. The `man` pages shows us that it takes 2 arguments -- the strings to compare.

Let's go find `rdi` and `rsi`.

```
0x00000000000001216 <+77>:    mov     rax,QWORD PTR [rbp-0x8]  
0x0000000000000121a <+81>:    lea     rdx,[rip+0xdf9]    # 0x201a  
0x00000000000001221 <+88>:    mov     rsi,rdx  
0x00000000000001224 <+91>:    mov     rdi,rax  
0x00000000000001227 <+94>:    call    0x10c0 <strcmp@plt>
```

From this, we notice two things:

- `rdi` is loaded with the address of `rbp-0x8`, which is where the `malloc` happened. This means that `rdi` is the first string.
- `rsi` is loaded with the address of `rip+0xdf9`, which is `0x201a`. This is not something that we loaded. We notice that it is based on the instruction pointer (because of PIE), which tells us that it is something hardcoded. *More on this later*, but for now let's check what's there.

```
gef> x/s 0x201a
0x201a: "flag{welcome_to_runtime}\n"
```

As we expected, this is our flag. The following puts statements aren't really relevant to us, but we can guess by the code there that it prints a "yes" or "no" based on the return value of `strcmp`.

Conclusions

Based on this program, we see that the flag was hardcoded. This is not going to be the case in 99% of the binary exploitation problems. There is going to be some reverse engineering problems where the flag is obfuscated and your challenge is to figure out how it is converted, but not for binary exploitation problems.

For your reference, here is the source code:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define STR_SIZE 0x20

int main(int argc, char* argv[])
{
    char* line = malloc(STR_SIZE * sizeof(char));
    printf("Enter the flag here: ");
    fgets(line, STR_SIZE, stdin);

    if (strcmp(line, "flag{welcome_to_runtime}\n") == 0)
        printf("That's the right flag!\n");
    else
        printf("Nope, try again!\n");

    return 0;
}
```

PS: You could have run `strings intro | grep flag` to find the flag. `strings` returns the hardcoded strings in binaries, so it's not a bad thing to check as a first step.